

The Elicitation Round

What to do, and why it is the last thing outstanding

This packet is everything needed to run the one open item on the project. Nothing in the repository is blocked on engineering. This is the only task left that could change a conclusion rather than qualify one.

Why it matters

Part Four of the book rests on a twelve-by-eight table that one person wrote down. Every robustness check in the project varies the *values* in that table and holds its pattern of empty cells fixed, so no amount of further computation can test the pattern itself. The book already says as much, and records that Part Four's coupling claims largely do not survive structural randomisation. A second reader filling the same grid independently is the only test available.

The result is genuinely two-sided. If two or three readers leave the same rows empty, a central claim has independent support for the first time. If they do not, **Chapter 18 is wrong**, and the preregistered script says so in those words rather than softening it.

The six documents

Document	Who sees it
0-start-here	you
1-respondent-form	the respondent — this is the document you send
2-recruiting-note	you, to copy into an email or message
3-response-template	send with the form, or let them return the form filled in
4-collator-notes	you only
5-invitation-email.md	you, to paste into an email

Do not send 4-collator-notes to a respondent. It states how many rows the book leaves empty and which they are, which is the answer the exercise is trying to elicit independently. Sending it destroys the round.

Steps

1. **Recruit two or three people.** Three is the useful minimum: with three you can report, per cell, whether zero, one, two or three marked it, and the empty-rows claim becomes a statement about agreement rather than about one person.

The ideal respondent knows AA well enough to have opinions about what the Traditions do, and has *not* read Part Four. Someone who has read it is not independent; that does not disqualify them, but it must be recorded, and their form reported separately.

2. **Send them 1-respondent-form** and the note in 2-recruiting-note, or the ready-to-send message in 5-invitation-email.md, which adds the subject lines, the attachment check, answers to the questions respondents actually ask, a nudge and a thank-you. Twenty to forty minutes of work for them.

Ask the respondents not to compare notes with each other until every form is back, and keep

your own replies to their questions short: the likeliest way to spoil the round is a helpful answer to a follow-up.

3. **Save each completed form** in the repository as `research/governance-elicitation-<initials>.md`. The script reads every file matching that pattern: a markdown table whose first column is the Tradition and whose remaining eight columns are the resources in the order given on the form. Blank and 0 both mean empty.

4. **Run the comparison:**

```
python3 model/elicitation_compare.py --self-test
python3 model/elicitation_compare.py
```

Run the self-test before sending any forms out. It passed on 10 August 2026.

5. **Record who filled each one in**, what they were told, and whether they had seen the book, in `research/SOURCES.md`. Then propagate whatever the result says through the layers listed in `CLAUDE.md`, exactly as any other change to a public claim.

The instrument was checked before this packet was made

`model/elicitation_compare.py --self-test` passed on 10 August 2026. It exercises two synthetic respondents: one who matches the book's empty rows, and one who differs by a single row. The second is correctly reported as contradicting Chapter 18 rather than smoothed over, which is the behaviour that matters, since an instrument that only confirms is not an instrument.

The analysis is preregistered. `model/elicitation_compare.py` was written before any form came back, so what gets reported cannot be chosen after seeing the answers.

What you cannot conclude from this

Three respondents is not a validation study. Agreement would show that the table is reproducible by people working from the same definitions, not that the definitions carve the world correctly. The resource list is itself one person's, and question 2 on the form invites a respondent to say it is wrong, which would be a more serious objection than any individual cell.